

**AdaCore**

Build Software that Matters

# **ParaSail: A Pointer-Free Pervasively-Parallel Language for Irregular Computations**

S. Tucker Taft, AdaCore

<Programming> 2019, Genoa, Italy

April 2019



# Goals for ParaSail

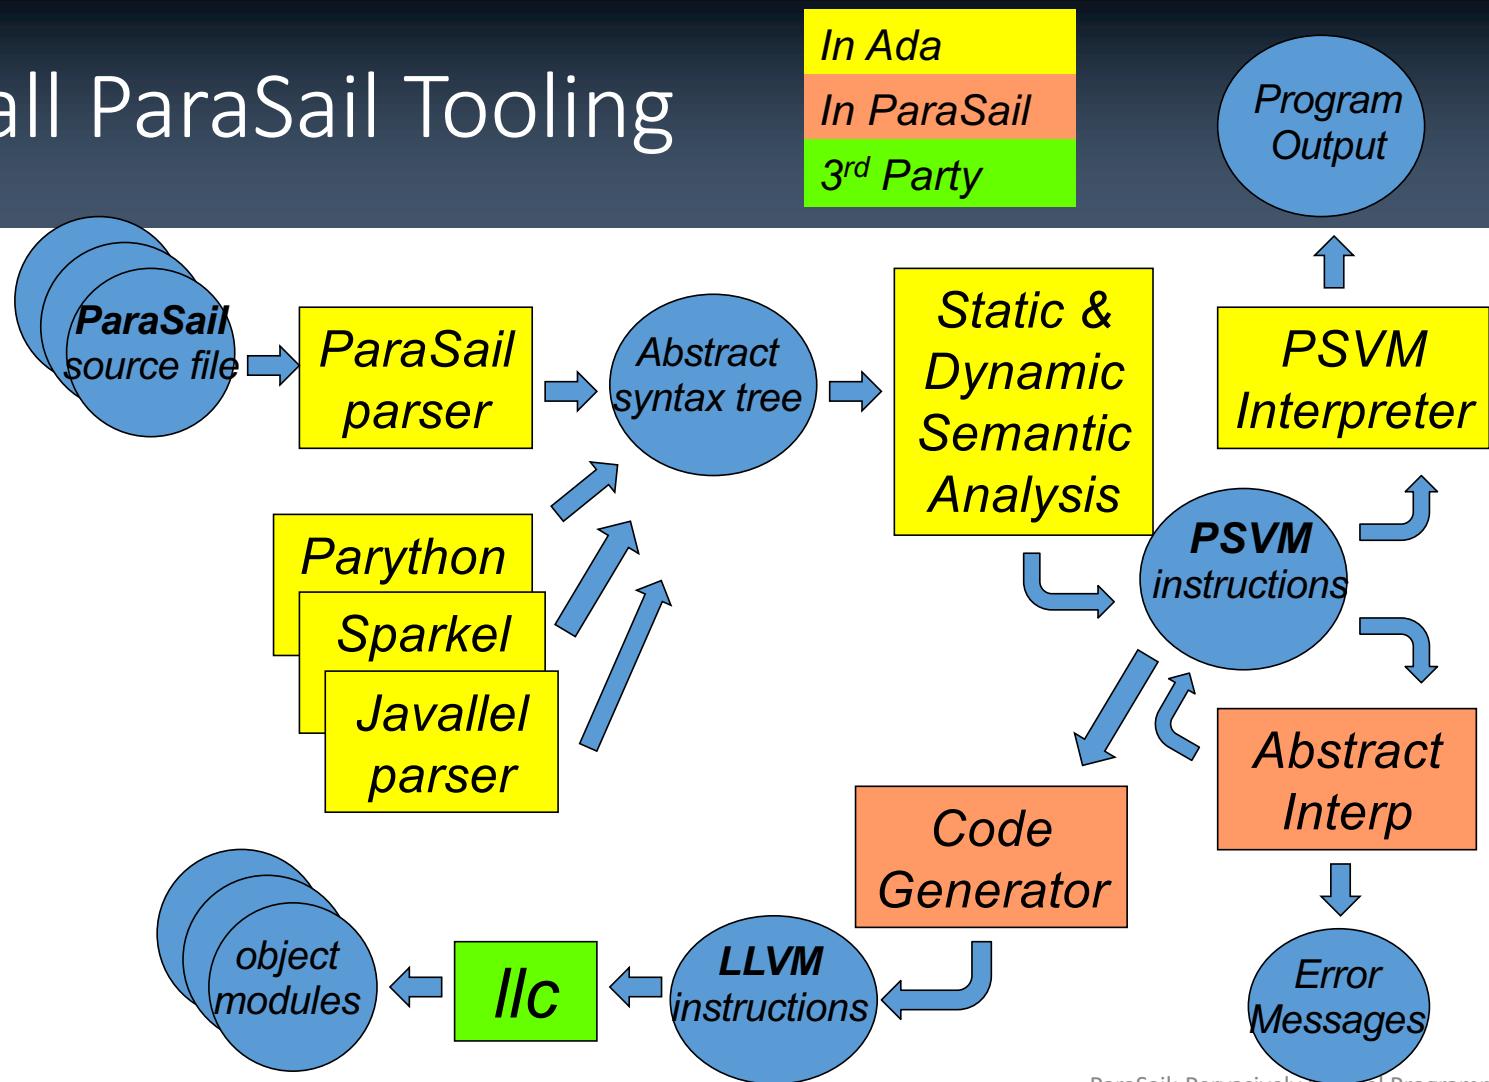
*Parallel Specification and implementation language*

- Goal A: Simplify writing parallel programs
  - Make it as easy for programmer to write in parallel as it is sequentially
  - Make it easy for compiler to detect data races
  - Make it easy for compiler to insert parallelism
- Goal B: Simplify writing safe, correct, and predictable code
  - Provide automatic storage management without garbage collection
  - Integrate contract-based programming annotations into language
  - Eliminate difficult-to-verify features

# History of ParaSail

- **2009 – Started design of ParaSail, on a blog**
  - My friends were too busy – drowned my sorrows in a **blog**
  - Design documented along the way on <http://parasail-programming-language.blogspot.com>
- **2011 – ParaSail interpreter completed**
- **2013 – Refactored to support multiple parsers, same AST**
  - *Javallel, Parython, Sparkel* – to address **surface syntax** preferences
- **2014 – LLVM-based code generator built**
  - Parallelism is only interesting if it makes the program **faster**
  - Written in interpreted **ParaSail** using *reflection* (by a summer intern)
- **2015 – Abstract-interpretation-based static analysis tool built:**
  - **ParaScope: ParaSail Static Catcher of Programming Errors**
  - Written in interpreted **ParaSail** using *reflection*
- **2017-8 – Refactored LLVM-based code generator**
  - Use **register** model rather than **stack** model for temps and parameters in LLVM
  - **Inline** more of the run-time support

# Overall ParaSail Tooling



# Start with Goal B: Simplify and Unify Language to “make room” for Parallelism and Contracts

- Modules and Types – Can we simplify and unify?
  - Package, Template, Namespace, Class, Interface, Object, Record, Struct, Structure, Sig, Signature, Monitor, Task
    - Do we really need more than one of these?
    - **What unifies these?**
  - Their components are heterogeneous at compile-time
  - A symbolic selector identifies a component and determines its type.
  - Some are parameterized, some are not.
  - Some are modules, some are types, some are both.
  - Some define an interface; some define an implementation with an implicit interface
- What about Operations – must we have so many kinds?
  - Virtual function, Static function, Constructor, Destructor, Method, Procedure
    - Is it Object.Operation(...) vs. Operation(Object, ...)?
    - **What unifies these?**
  - Algorithm with inputs, outputs, and local variables, and possibly up-level access to globals
  - Inheritance of interface and optionally implementation across type hierarchy
    - Indirect binding as part of interface inheritance
- Objects – value vs. reference model? mutable vs. immutable?

# ParaSail “Simplified” Building Blocks



- **Module** – *always parameterized, with separate interface*
  - Parameterized unit with *Interface* part, and optionally an *Implementation* part
  - **interface** Enum <Literals : Vector <Univ\_Enumeration>> **is** ...
  - Can *extend* another module (code and data are inherited along with interface)
  - Can *implement* one or more module's interfaces (only interface is inherited)
- **Type** – *single syntax for all type declarations*
  - An *instance* of a Module: **type** T **is** [new] M <...> // “new” means use “name” equivalence
  - e.g. **type** Color **is new** Enum <[#red, #green, #blue]>
  - “T+” is *polymorphic* type representing any type implementing T’s interface
- **Object** – *only “value” types, var or const*
  - An *instance* of a Type (type can be inferred from initial value)
    - has an updatable *value* if declared “**var**”; can be **null** if declared “**optional**”
  - **var** Obj : T := Create(...) // ParaSail allows overloading on parameter and result types
- **Operation** – *role of operation determined by parameters and where declared*
  - *Defined* in a Module, and
  - *Operates* on one or more Objects of specified Types
  - Only operates on its *explicit* parameters; can *update* only those declared “**var**”

# A Map Module Interface in ParaSail

```
interface PSL::Containers::Map           // A hashed-map module; a set of K/V Pairs
  <Key_Type is Hashable<>; Value_Type is Assignable<>> // formal params of module
is
  op "[]"() -> Map                  // Return an empty map
  type Pair is Key_Value<Key_Type, Value_Type> // K/V pair -- actual params in <...>
  op "|="(var M: Map; KV: Pair)      // Add [Key=>Value] pair to Map
  op "|"(M: Map; KV: Pair) -> Map    // Return Map with [Key=>Value] pair added
  op "index_set"(M: Map) -> Set<Key_Type> // Return set of keys of the map
  op "in"(Key: Key_Type; M: Map) -> Boolean // True if Key in "index_set"(M)
  op "-="(var M: Map; Key: Key_Type)     // Remove mapping for Key, if any
  op "indexing"(ref M: Map; Key: Key_Type) {Key in M} -> ref Value_Type
    // Used for indexing via "M[Key]"; {...} is precondition if before "->"
  func Remove_Any(var M: Map) -> optional Pair
    // Remove one mapping from Map; Return null if map is empty.
  op "magnitude"(M: Map) -> Univ_Integer
    // Number of mappings in the table; Supports "|M|" notation
  func Is_Empty(M: Map) -> Boolean // True if no mappings in M
end interface PSL::Containers::Map // end bracketing is optional if indented properly
```

## Key to listing:

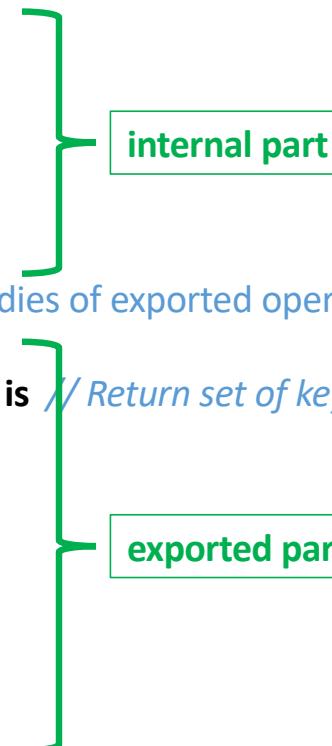
**func**: callable function/method/operation  
**op**: function that relies on *syntactic sugar*  
**var**: allows update of parameter  
**ref**: short-lived reference to existing object  
**type**: nested type; instance of module  
**{...}**: contract (pre/postcondition, invariant, ...)

# Example uses of Map Module

<b>var</b> X : Map <Key_Type => String, Value_Type => Integer> := []	<i>mutable map, string =&gt; integer</i>
<b>const</b> Y : Map <String, Integer> := [“the” => 35, “cat” => 42]	<i>immutable map; uses “[]” &amp; “ =”</i>
X  = [“dog” => 2]	<i>add key/val pair to X</i>
<b>var</b> Z := Y   X	<i>type inferred from initial value</i>
<b>type</b> String_Int_Map is Map<String, Value_Type => Integer>	<i>named type, but equiv to anon</i>
<b>type</b> Word_Counts is new Map<String, Integer>	<i>named type, not equiv to anon</i>
<b>var</b> WCs : Word_Counts := Z // illegal – type mismatch	<i>“new” forces name equivalence</i>
<b>for each</b> [W => C] <b>of</b> Z <b>loop</b> <i>// could be {forward, reverse, concurrent} loop</i> IO.Stdout.Println(“Word: ` (W), Count: ` (C)”) <b>end loop</b> // end bracketing is optional	<i>syntactic sugar uses operations of Map (“index_set”, “indexing”) and Set (Remove_Any [or Remove_First, Remove_Last])` (...) is string interpolation syntax</i>

# Parts of Map Module Implementation

```
class Map is
    interface Node<> is // nested local module
        var Entry: Pair
        var Next: optional Node
    end interface Node
    var Backbone: Basic_Array<optional Node>
    var Count: Univ_Integer := 0
    exports // "exports" separates internal part from bodies of exported operations
    ...
    op "index_set"(M : Map) -> Result : Set<Key_Type> is // Return set of keys of the map
        Result := []
        for each B of T.Backbone loop
            for N => B then N.Next while N not null loop
                Result |= Key_Of (N.Entry)
            end loop
        end loop
    end op "index_set"
    ...
end class Map //end bracketing optional
```



# Aside: How to live without pointers? Use *Expandable* Objects Instead

- All types have additional **null** value, which takes very little space
  - Only objects or components declared **optional** can be null; allows object to grow and shrink
  - Eliminates many of the common uses for pointers, e.g. trees
  - Assignment (“`:=`”) is by copy
    - Move (“`<==`”) and swap (“`<=>`”) operators also provided
    - Essentially equivalent to an implicit “ownership” model (cf. Rust, `unique_ptr`, etc.)
- Generalized indexing into objects replaces pointers for cyclic graph structures
  - **for each N of** `Directed_Graph[I].Successors` **loop** ...
- Region-Based Storage Mgmt can replace Global Heap
  - All objects are “local” with growth/shrinkage using local heap
  - “**null**” value carries indication of region to use on growth
- Short-lived references to existing objects are permitted
  - Returned by user-defined indexing functions, for example
  - Useful for defining “slices” of an array or other container

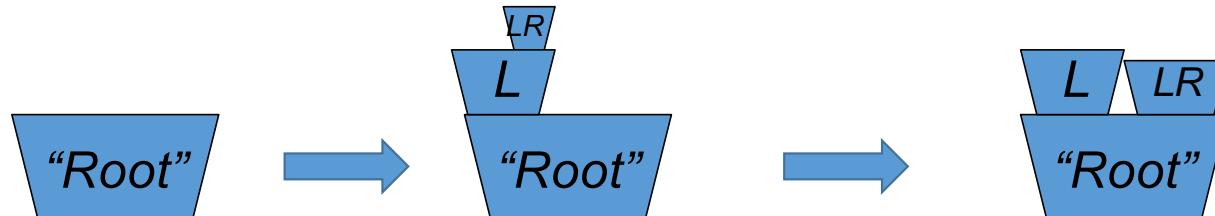
# Example: Pointer-Free Trees

```
interface Tree_Node <Contents_Type is Assignable>> is
    var Contents: Contents_Type
    var Left : optional Tree_Node := null
    var Right : optional Tree_Node := null
end interface Tree_Node
```

```
var Root : Tree_Node<Univ_String> := (Contents => "Root")
Root.Left := (Contents => "L", Right => (Contents => "LR"))
Root.Right <== Root.Left.Right // Root.Left.Right now null
```



*The Flower-Pot model of objects*

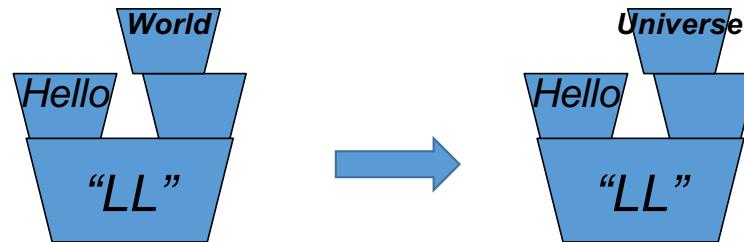


# Example: Pointer-Free Linked List

```
interface Linked_List<Elem_Type> is
    var Data : Elem_Type
    var Next : optional Linked_List := null
end interface Linked_List
```

```
LL : Linked_List <String> :=
    (Data => "Hello", Next => (Data => "World"))
```

```
LL.Next.Data := "Universe"
```



# And now Goal A: Pervasive Parallel Programming

- Make parallel evaluation the default
  - $F(X) + G(Y)$ 
    - Goal: safely evaluate  $F(X)$  and  $G(Y)$  in parallel without looking inside  $F$  &  $G$
    - Even if  $F$  has side effects on  $X$ , or  $G$  has side effects on  $Y$
  - Seemingly trivial with pure functional programming language
    - But would disallow any update-in-place optimizations
    - Monads can re-create their own parallelization challenges
    - Any “breaking” of purity as in ML family “ref”s brings back all of the challenges
  - Hard if  $X$  and  $Y$  use unrestricted pointers, since both might point to “ $Z$ ”
  - Hard if  $F$  and  $G$  can make unrestricted use of globals w/o explicit synchronization
    - But any synchronization would be bad for performance

# Goals A & B: Eliminate features that get in the way

- Eliminate global variables
  - Operation can only access or update variable state via its parameters
- Eliminate parameter aliasing
  - Use “hand-off” semantics (cf. Hermes language)
- Eliminate explicit threads, lock/unlock, signal/wait
  - Concurrent objects synchronized automatically
- Eliminate run-time exception handling
  - Can use compile-time checking and propagation of preconditions
  - “Mutually Assured Destruction” of sibling threads on exit is supported as an alternative.
- Eliminate pointers
  - Adopt notion of “optional” objects that can grow and shrink
  - User-defined indexing as alternative for cyclic graph structures
- Eliminate global heap and garbage collector
  - Replaced by region-based storage management (local heaps)
  - All objects conceptually live in a local stack frame

# Example: An Implicitly Parallel Divide-and-Conquer Recursive Not-In-Place (“functional”) QSort

```
func QSort(V : Vec_Type is Vector<Comparable<>>) -> Vec_Type is
    if |V| <= 1 then
        return V                                // The easy case
    else
        const Mid := V[ |V|/2 ] // Pick a pivot value
        return // implicit parallelism inserted here on recursive calls
            QSort( [for each E of V {E < Mid} => E] ) // Recurse; {. . .} is filter
            | [for each E of V {E == Mid} => E]      // No recursion since all values equal the pivot
            | QSort( [for each E of V {E > Mid} => E] ) // Recurse
    end if
end func QSort
```

Number of Cores	Core Configuration	Wall Clock	CPU	Thread Utilization	Speedup
1	Single Threaded	77.0	76.0	98.6	1 x
2 x 2	Dual Core x 2 Hyper Threads/Core	42.6	142.2	333.8	1.8 x

# An Explicitly Parallel In-Place Non-Recursive Qsort

(see paper for additional explication)

```
func Quicksort(var A : Array_Type) is // sort array A in place
    for Arr => A[..] while |Arr| > 1 loop // process slices
        if |Arr| == 2 then // short slice: simple case
            if Arr[Arr.Last] < Arr[Arr.First] then
                Arr[Arr.First] <=> Arr[Arr.Last] // swap
        else // long slice: partition array
            const Mid := Arr[Arr.First + |Arr|/2]
            var Left := Arr.First
            var Right := Arr.Last
        *Swapping* until Left > Right loop // keep swapping
            var New_Left := Right+1
            var New_Right := Left-1
            then // "then" waits for prior processing
            for I in Left .. Right forward loop // scan from left
                if Mid <= Arr[I] then
                    New_Left := I // Item OK for right partition
                if Mid < Arr[I] then
                    exit loop // Item must go in right partition
            || // "||" requests concurrent processing
```

```
for J in Left .. Right reverse loop // scan from right
    if Arr[j] <= Mid < Arr[J] then
        New_Right := J // Item OK for left partition
    if Arr[J] < Mid then
        exit loop // Item must go in left partition
    then // waits for the above "for I ... || for J ..." to complete
    if New_Left > New_Right then // check if any more swaps
        Left := New_Left
        Right := New_Right
        exit loop Swapping // no more swapping needed
    Arr[New_Left] <=> Arr[New_Right] // swap items
    Left := New_Left + 1 // skip past swapped items
    Right := New_Right - 1
end loop Swapping // continue looking for items to swap
then // add new slices to "bag" of slices to sort in parallel
    continue loop with Arr => Arr[Arr.First .. Right]
|| // "||" is higher precedence than "then"
    continue loop with Arr => Arr[Left .. Arr.Last]
```

# An interesting Parallel programming idiom

Concurrent loop(s) plus explicit “**continue loop with X => ...**”

Replaces recursion with parallel worklist-like idiom

*Example 1: Qsort of A:*

```
for Arr => A[..] while |Arr| > 1 loop // start with full-array "slice" A[..]
    ... partition slice "Arr" using swapping, and process two sub-slices
    // use "continue loop" to add work items to the "bag" of iterations
    continue loop with Arr => Arr[Arr.First .. Right]
    || // "||" is pronounced "in parallel with ..."
    continue loop with Arr => Arr[Left .. Arr.Last]
end loop // wait for "bag" of iterations to complete
```

# “Bag of Iterations” Idiom, Example 2: N-Queens

```
const Rows : Countable_Range<Chess_Unit> := 1..N
var Solutions : concurrent Vector<Board_State> := [] // “concurrent” means locking

*Outer_Loop*
for Partial : Board_State := No_Columns(N) loop      // Start with no columns filled
    for R in Rows concurrent loop                    // Try all rows in next column
        const Next := Add_Queen(Partial, At_Row => R) // Add queen at given row
        if Is_Safe(Next) then                          // See if is a safe place for queen
            if Num_Columns(Partial) == N then          // New queen is OK; see if done
                Solutions |= Next                      // Remember another complete solution
            else
                continue loop Outer_Loop
                with Partial => Next // Add partial solution to bag of (Outer) work items
            end if
        end if
    end loop
end loop Outer_Loop
```

# Other ParaSail Tidbits

- Compile-time computation thanks to full-language interpreter
  - Can use types produced by module instantiations
  - Literals, any self-contained computation, computed at compile-time
    - Can use assertions/pre-conditions for “type checking”
- Heavy use of Syntactic sugar; Syntax agnostic semantics
  - Simple unified underlying model for defining abstractions
    - No “implicit” parameter in decl, but allow Obj.Operation(...) notation
  - Compact/familiar/expressive syntax for use of abstractions by client
  - Supports multiple surface syntaxes
- Safe Distributed and Concurrent Programming
  - pointer-free (logical) objects are “self contained”
- Contract Annotations in the language
  - recognized by both the interpreter and the compiler
  - compile-time proved, run-time checked, or hybrid

# Evaluation of ParaSail against its Goals

- *Use syntax and type model familiar to existing professional programmers*
  - Anecdotal evidence for readability and understandability from existing C++, Java, Python programmers
    - Though a bit verbose for some!
  - Support for alternative surface syntax is of interest to some
- *Ease the creation of safe, parallel programs of significant size*
  - LLVM code generator written by parallel programming neophyte during 6-month internship
  - High levels of parallelism achieved without significant extra effort
    - Lack of globals and compile-time race detection biggest help
- *Raise the level of abstraction*
  - Static analysis tool based on abstract interpretation re-implemented in ParaSail
  - Value Propagation algorithm significantly simpler despite lack of pointers and globals
    - Largely thanks to pointer-free data structuring and concise syntactic sugar available to all user-defined data types
- *Support real-time embedded parallel programming*
  - Support for both compile-time and run-time synchronization
  - Bounded, predictable storage management
- *Achieve these goals with economy of means*
  - Simple Unified Module/Type/Object/Operation Building Blocks
  - “50-page” reference manual a la Modula-3

# Next Steps for ParaSail

- *Performance, performance, performance*
  - Inlining more of the run-time support
  - Stack-resident (**non-optional** or bounded) compound objects
  - Inlining **non-optional** component objects into enclosing object
  - Eliminating header overhead for **non-optional** component objects
- *Parython*
  - Productize Python variant of ParaSail
  - Parallel, type-checked, interpreted or compiled-to-LLVM Python
- *Take more advantage of full interpreter at compile-time*
  - Currently used for checking properties of literals
    - Enum Literal is part of enumeration, Integer literal is in range, String literal has desired form
  - Any ParaSail expression with no use of variables is evaluated at compile time
    - Could formalize into a separate “user-directed type-check” step
      - cf. Ruby **comp-types** by Jeff Foster (Tufts/UMD)

# Conclusion : Overall ParaSail Themes

- *Mutable Objects* with Value Semantics
- *Stack-Based* Heap Management
- *Compile-Time* Exception Handling
- *Race-Free* Parallel Programming

Expressivity of functional programming, but with safe mutation, easy parallelization, and ability to model the “real” world

*Downloadable from: [parasail-lang.org](http://parasail-lang.org)*

*Forum: [groups.google.com/group/parasail-programming-language/](http://groups.google.com/group/parasail-programming-language/)*

*Contact: [taft@adacore.com](mailto:taft@adacore.com)*

*Blog: <http://parasail-programming-language.blogspot.com/>*